

1.0 BUSINESS UNDERSTANDING

1.1 Background

SyriaTel is one of the leading telecommunication companies in Syria. It provides a wide range of telecommunications services, including mobile and fixed-line telephony, internet services and data services.

Syriatel has been a key player in the Syrian telecommunications market, serving millions of customers across the country. The company has played a significant role in expanding and modernizing telecommunications infrastructure, contributing to the country's connectivity and economic development.

The telecommunications company is interested in reducing how much money is lost because of customers who don't stick around very long.

1.2 Problem Statement

In this competitive world, business is becoming highly saturated. Especially, the field of telecommunication faces complex challenges due to a number of vibrant competitive service providers. Therefore, it has become very difficult for them to retain existing customers. Since the cost of acquiring new customers is much higher than the cost of retaining the existing customers, it is the time for the telecom industries to take necessary steps to retain the customers to stabilize their market value.

1.3 Objectives

1. To calculate the churn rate at SyriaTel, a telecommunications company.
2. To identify the factors that lead to churn and those that help in customer retention.

2.0 DATA UNDERSTANDING

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
import numpy as np
import re
import warnings
warnings.filterwarnings("ignore")
from sklearn.preprocessing import LabelEncoder, OneHotEncoder,
StandardScaler, normalize
from sklearn.model_selection import train_test_split, cross_val_score,
RandomizedSearchCV
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.tree import DecisionTreeClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import GaussianNB
from sklearn.ensemble import RandomForestClassifier,
```

```

GradientBoostingClassifier
from sklearn.svm import SVC
# %pip install xgboost
from xgboost import XGBClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import confusion_matrix, accuracy_score

```

```

df=pd.read_csv(r"C:\Users\user\Desktop\Phase_3_project\
customer_churn_dataset.csv")
df.head(5)

```

	state	account length	area code	phone number	international plan \
0	KS	128	415	382-4657	no
1	OH	107	415	371-7191	no
2	NJ	137	415	358-1921	no
3	OH	84	408	375-9999	yes
4	OK	75	415	330-6626	yes

	voice mail plan number	vmail messages	total day minutes	total day calls \
0	yes	25	265.1	110
1	yes	26	161.6	123
2	no	0	243.4	114
3	no	0	299.4	71
4	no	0	166.7	113

	total day charge	...	total eve calls	total eve charge \
0	45.07	...	99	16.78
1	27.47	...	103	16.62
2	41.38	...	110	10.30
3	50.90	...	88	5.26
4	28.34	...	122	12.61

	total night minutes	total night calls	total night charge \
0	244.7	91	11.01
1	254.4	103	11.45
2	162.6	104	7.32
3	196.9	89	8.86
4	186.9	121	8.41

	total intl minutes	total intl calls	total intl charge \
0	10.0	3	2.70
1	13.7	3	3.70
2	12.2	5	3.29

3	6.6	7	1.78
4	10.1	3	2.73

	customer service calls	churn
0	1	False
1	1	False
2	0	False
3	2	False
4	3	False

[5 rows x 21 columns]

Better view of the dataframe since it has many columns
Transposing does not transform or modify the original df. It only
enhances the visibility of the many columns

df.head().T

	0	1	2	3
4				
state	KS	OH	NJ	OH
OK				
account length	128	107	137	84
75				
area code	415	415	415	408
415				
phone number	382-4657	371-7191	358-1921	375-9999
6626				
international plan	no	no	no	yes
yes				
voice mail plan	yes	yes	no	no
no				
number vmail messages	25	26	0	0
0				
total day minutes	265.1	161.6	243.4	299.4
166.7				
total day calls	110	123	114	71
113				
total day charge	45.07	27.47	41.38	50.9
28.34				
total eve minutes	197.4	195.5	121.2	61.9
148.3				
total eve calls	99	103	110	88
122				
total eve charge	16.78	16.62	10.3	5.26
12.61				
total night minutes	244.7	254.4	162.6	196.9
186.9				
total night calls	91	103	104	89
121				

total night charge	11.01	11.45	7.32	8.86
8.41				
total intl minutes	10.0	13.7	12.2	6.6
10.1				
total intl calls	3	3	5	7
3				
total intl charge	2.7	3.7	3.29	1.78
2.73				
customer service calls	1	1	0	2
3				
churn	False	False	False	False
False				

number of rows and columns in the dataframe
Each row represents a record of a customer

```
print("The number of rows in the SyriaTel dataframe is", df.shape[0])
print("The number of columns in the SyriaTel dataframe is",
df.shape[1])
```

The number of rows in the SyriaTel dataframe is 3333
The number of columns in the SyriaTel dataframe is 21

Getting column information such as the Datatype and number of non-null values
df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3333 entries, 0 to 3332
Data columns (total 21 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   state                                3333 non-null   object
1   account length                       3333 non-null   int64
2   area code                           3333 non-null   int64
3   phone number                         3333 non-null   object
4   international plan                   3333 non-null   object
5   voice mail plan                      3333 non-null   object
6   number vmail messages                3333 non-null   int64
7   total day minutes                    3333 non-null   float64
8   total day calls                      3333 non-null   int64
9   total day charge                     3333 non-null   float64
10  total eve minutes                    3333 non-null   float64
11  total eve calls                      3333 non-null   int64
12  total eve charge                     3333 non-null   float64
13  total night minutes                  3333 non-null   float64
14  total night calls                    3333 non-null   int64
15  total night charge                   3333 non-null   float64
16  total intl minutes                   3333 non-null   float64
17  total intl calls                     3333 non-null   int64
```

```

18 total intl charge      3333 non-null    float64
19 customer service calls 3333 non-null    int64
20 churn                  3333 non-null    bool
dtypes: bool(1), float64(8), int64(8), object(4)
memory usage: 524.2+ KB

```

Getting the names of the columns

```
df.columns
```

```

Index(['state', 'account length', 'area code', 'phone number',
      'international plan', 'voice mail plan', 'number vmail
messages',
      'total day minutes', 'total day calls', 'total day charge',
      'total eve minutes', 'total eve calls', 'total eve charge',
      'total night minutes', 'total night calls', 'total night
charge',
      'total intl minutes', 'total intl calls', 'total intl charge',
      'customer service calls', 'churn'],
      dtype='object')

```

getting the descriptive statistics of the dataframe

```
df.describe()
```

	account length	area code	number vmail messages	total day minutes \
count	3333.000000	3333.000000	3333.000000	3333.000000
mean	101.064806	437.182418	8.099010	179.775098
std	39.822106	42.371290	13.688365	54.467389
min	1.000000	408.000000	0.000000	0.000000
25%	74.000000	408.000000	0.000000	143.700000
50%	101.000000	415.000000	0.000000	179.400000
75%	127.000000	510.000000	20.000000	216.400000
max	243.000000	510.000000	51.000000	350.800000

	total day calls	total day charge	total eve minutes	total eve calls \
count	3333.000000	3333.000000	3333.000000	3333.000000
mean	100.435644	30.562307	200.980348	100.114311

std	20.069084	9.259435	50.713844
19.922625			
min	0.000000	0.000000	0.000000
0.000000			
25%	87.000000	24.430000	166.600000
87.000000			
50%	101.000000	30.500000	201.400000
100.000000			
75%	114.000000	36.790000	235.300000
114.000000			
max	165.000000	59.640000	363.700000
170.000000			

	total eve charge	total night minutes	total night calls \
count	3333.000000	3333.000000	3333.000000
mean	17.083540	200.872037	100.107711
std	4.310668	50.573847	19.568609
min	0.000000	23.200000	33.000000
25%	14.160000	167.000000	87.000000
50%	17.120000	201.200000	100.000000
75%	20.000000	235.300000	113.000000
max	30.910000	395.000000	175.000000

	total night charge	total intl minutes	total intl calls \
count	3333.000000	3333.000000	3333.000000
mean	9.039325	10.237294	4.479448
std	2.275873	2.791840	2.461214
min	1.040000	0.000000	0.000000
25%	7.520000	8.500000	3.000000
50%	9.050000	10.300000	4.000000
75%	10.590000	12.100000	6.000000
max	17.770000	20.000000	20.000000

	total intl charge	customer service calls
count	3333.000000	3333.000000
mean	2.764581	1.562856
std	0.753773	1.315491
min	0.000000	0.000000
25%	2.300000	1.000000
50%	2.780000	1.000000
75%	3.270000	2.000000
max	5.400000	9.000000

Observations

The minimum number of voicemail messages, total day minutes, total day calls, total day charge, total evening minutes, total evening calls and total evening charge are all 0.

```
# Understanding the current Churn Status
```

```
df["churn"].value_counts()
```

```
False    2850  
True      483  
Name: churn, dtype: int64
```

It is evident that there are 483 disloyal customers who have churned the services of SyriaTel

3.0 DATA PREPARATION.

3.1 Data Cleaning

```
# checking for duplicates
```

```
df.duplicated().value_counts()
```

```
False    3333  
dtype: int64
```

Observation: There are no duplicates in the 3333 records.

```
# Checking for missing values
```

```
df.isna().sum()
```

```
state                0  
account length      0  
area code           0  
phone number        0  
international plan   0  
voice mail plan      0  
number vmail messages 0  
total day minutes    0  
total day calls      0  
total day charge     0  
total eve minutes    0  
total eve calls      0  
total eve charge     0  
total night minutes  0  
total night calls    0  
total night charge   0  
total intl minutes   0  
total intl calls     0  
total intl charge    0  
customer service calls 0  
churn                0  
dtype: int64
```

Observation: No column has a missing value

```
# Checking the unique values of the states
df.state.unique()
array(['KS', 'OH', 'NJ', 'OK', 'AL', 'MA', 'MO', 'LA', 'WV', 'IN',
      'RI',
      'IA', 'MT', 'NY', 'ID', 'VT', 'VA', 'TX', 'FL', 'CO', 'AZ',
      'SC',
      'NE', 'WY', 'HI', 'IL', 'NH', 'GA', 'AK', 'MD', 'AR', 'WI',
      'OR',
      'MI', 'DE', 'UT', 'CA', 'MN', 'SD', 'NC', 'WA', 'NM', 'NV',
      'DC',
      'KY', 'ME', 'MS', 'TN', 'PA', 'CT', 'ND'], dtype=object)

# checking for the number of unique values in state
no_of_unique=df.state.nunique()

print("Observation: There are {} unique states represented in this
Dataframe.".format(no_of_unique))

Observation: There are 51 unique states represented in this Dataframe.

# Removing phone number as it has no use here
df.drop(columns=["phone number"], inplace = True)

# Confirming that we are now working with 20 columns

print("Number of columns are
now",df.columns.value_counts().sum(),"after removing phone number
column.")

Number of columns are now 20 after removing phone number column.
```

3.2 EXPLORATORY DATA ANALYSIS

3.2.1 Univariate Analysis

```
sns.countplot(x="churn", data=df)
plt.title ("Customer Churn Rate")

# Calculating churn percentages
total_count = df.shape[0] #rows/total num of customers
churn_counts = df["churn"].value_counts() # value count of each of
the two unique values
churn_percentages = churn_counts / total_count * 100

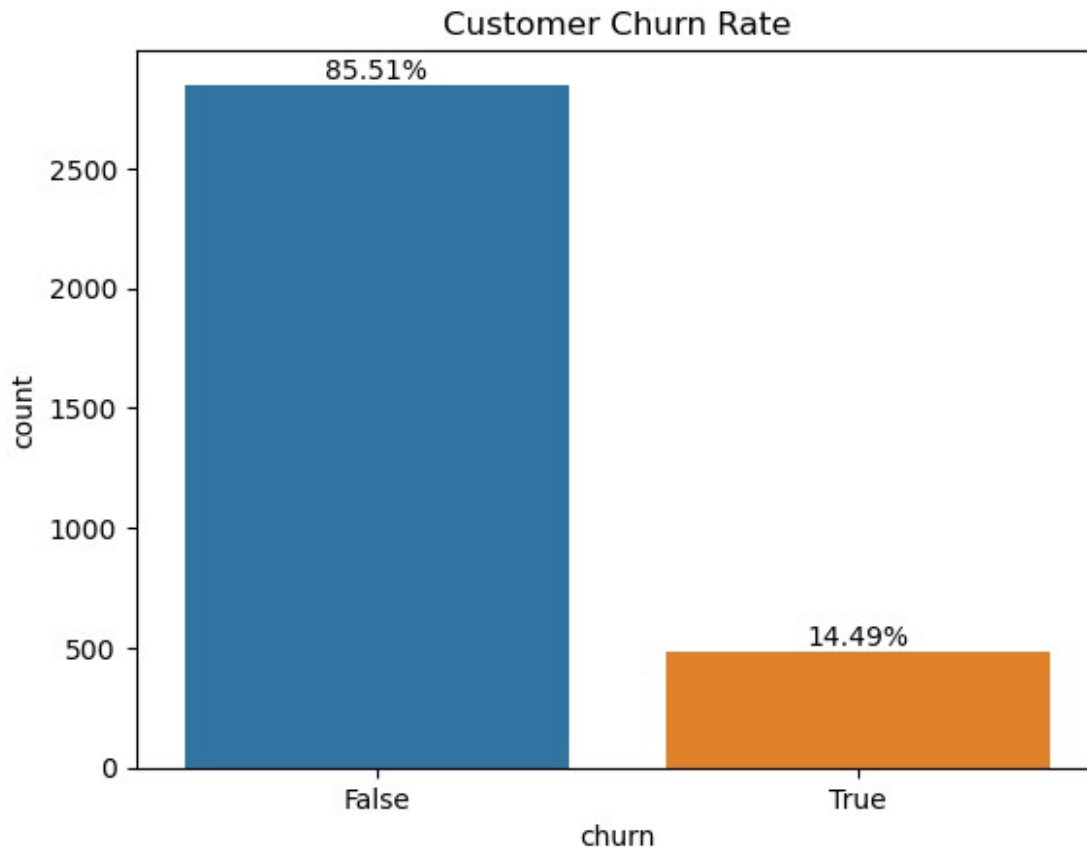
# Annotating the bars with churn percentages
```



```

for i, percentage in enumerate(churn_percentages):
    plt.text(i, churn_counts[i], f'{percentage:.2f}%', ha="center",
va= "bottom")
plt.show()

```



Observation: The churn rate currently stands at 14.49%

```

# State column will be plotted separately from the countplot,
# despite being nominal. This will be done due to high number of
unique states(51) which cannot fit well in a countplot.

columns_to_plot = [col for col in df.columns if col != "state"]

# Setting up the figure and axes for subplots
fig, axes = plt.subplots(nrows=5, ncols=4, figsize=(20, 20))

# Flattening the axes for easier iteration
axes = axes.flatten()

# Iterating over each column and create countplot(nominal/object) or
histogram(discreet)

```

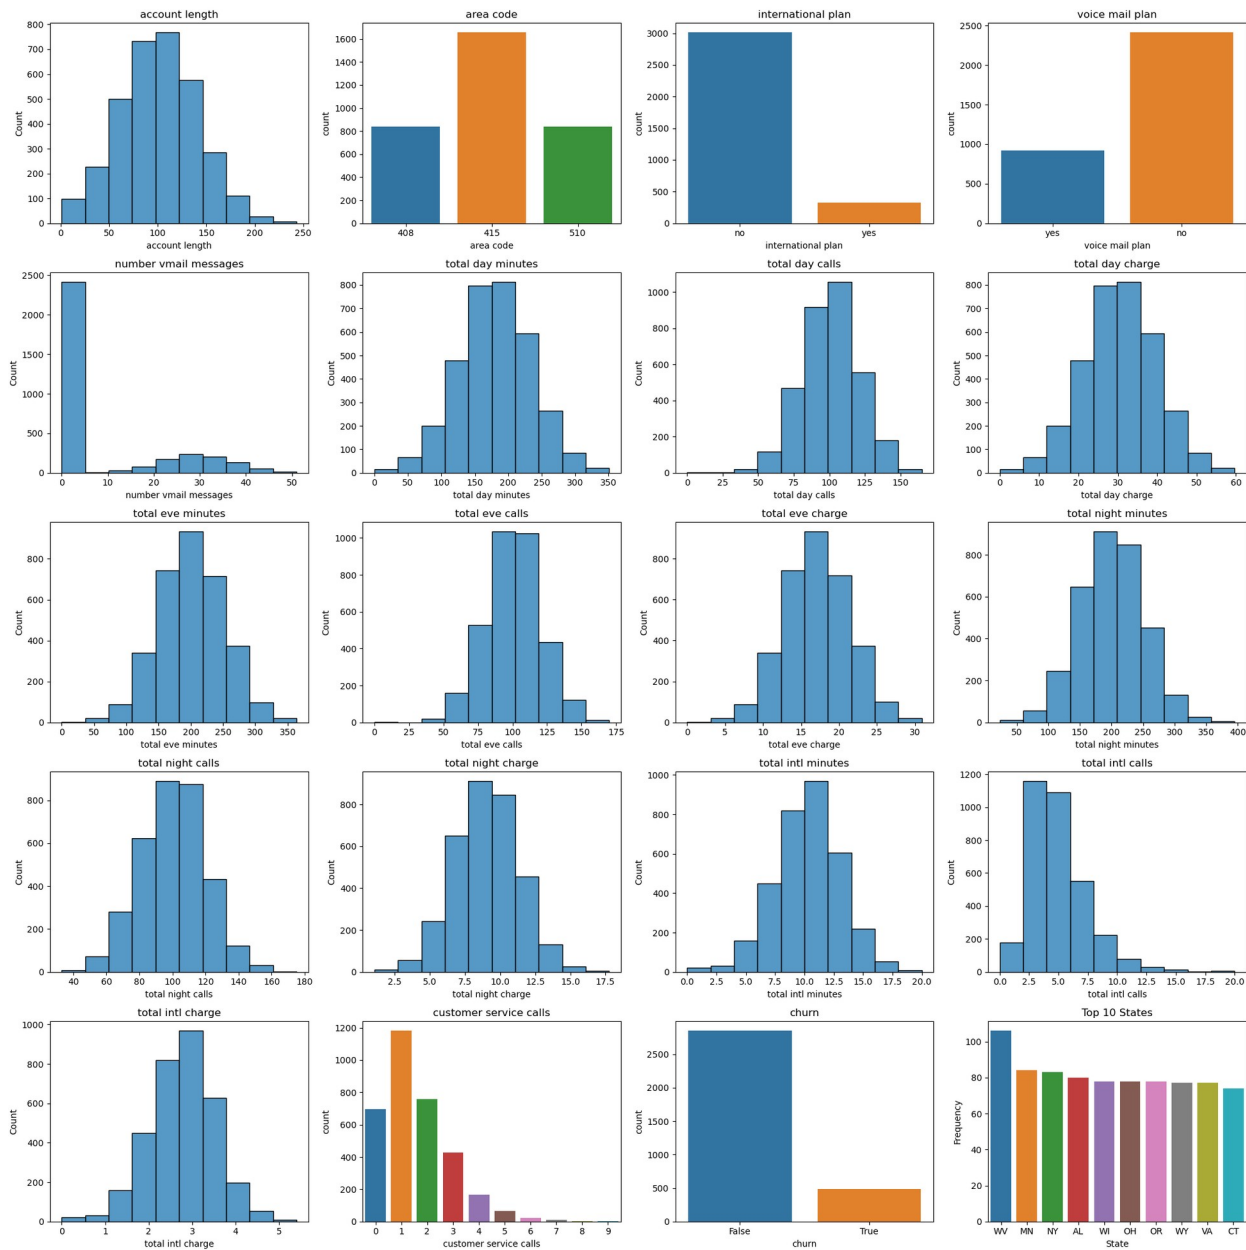
```
for i, column in enumerate(columns_to_plot):
    if df[column].dtype == 'object' or df[column].nunique() < 20:
        sns.countplot(x=column, data=df, ax=axes[i])
    else:
        sns.histplot(df[column], bins= 10, kde=False, ax=axes[i])
    axes[i].set_title(column)  # title for each subplot

# now i will plot the top 10 states

state_counts = df['state'].value_counts().head(10) # Top 10 states
sns.barplot(x=state_counts.index, y=state_counts.values)
plt.title('Top 10 States')
plt.xlabel('State')
plt.ylabel('Frequency')

# Adjusting layout to prevent overlap
plt.tight_layout()

plt.show()
```



```
skewness=df.skew()
skewness
```

```
account length      0.096606
area code           1.126823
number vmail messages 1.264824
total day minutes   -0.029077
total day calls     -0.111787
total day charge    -0.029083
total eve minutes   -0.023877
total eve calls     -0.055563
total eve charge    -0.023858
```

```
total night minutes    0.008921
total night calls      0.032500
total night charge     0.008886
total intl minutes     -0.245136
total intl calls       1.321478
total intl charge      -0.245287
customer service calls 1.091359
churn                  2.018356
dtype: float64
```

Observation:

Based on the distribution and information above, The following columns were highly skewed (values > 1 or <-1):

- 1.Area code The distribution is highly skewed to the right. There are likely a few area codes that occur much more frequently than others.
- 2.Number of voicemail messages The distribution is highly skewed to the right. Most customers likely have few or no voicemail messages, with a few customers having many.
- 3.Total international calls The distribution is highly skewed to the right. Most customers likely make very few international calls, with a few customers making many.
- 4.Customer service calls The distribution is highly skewed to the right. Most customers likely make very few customer service calls, with a few customers making many.
- 5.Churn The target variable was highly skewed to the right. This indicates that most customers did not churn, but a few did.

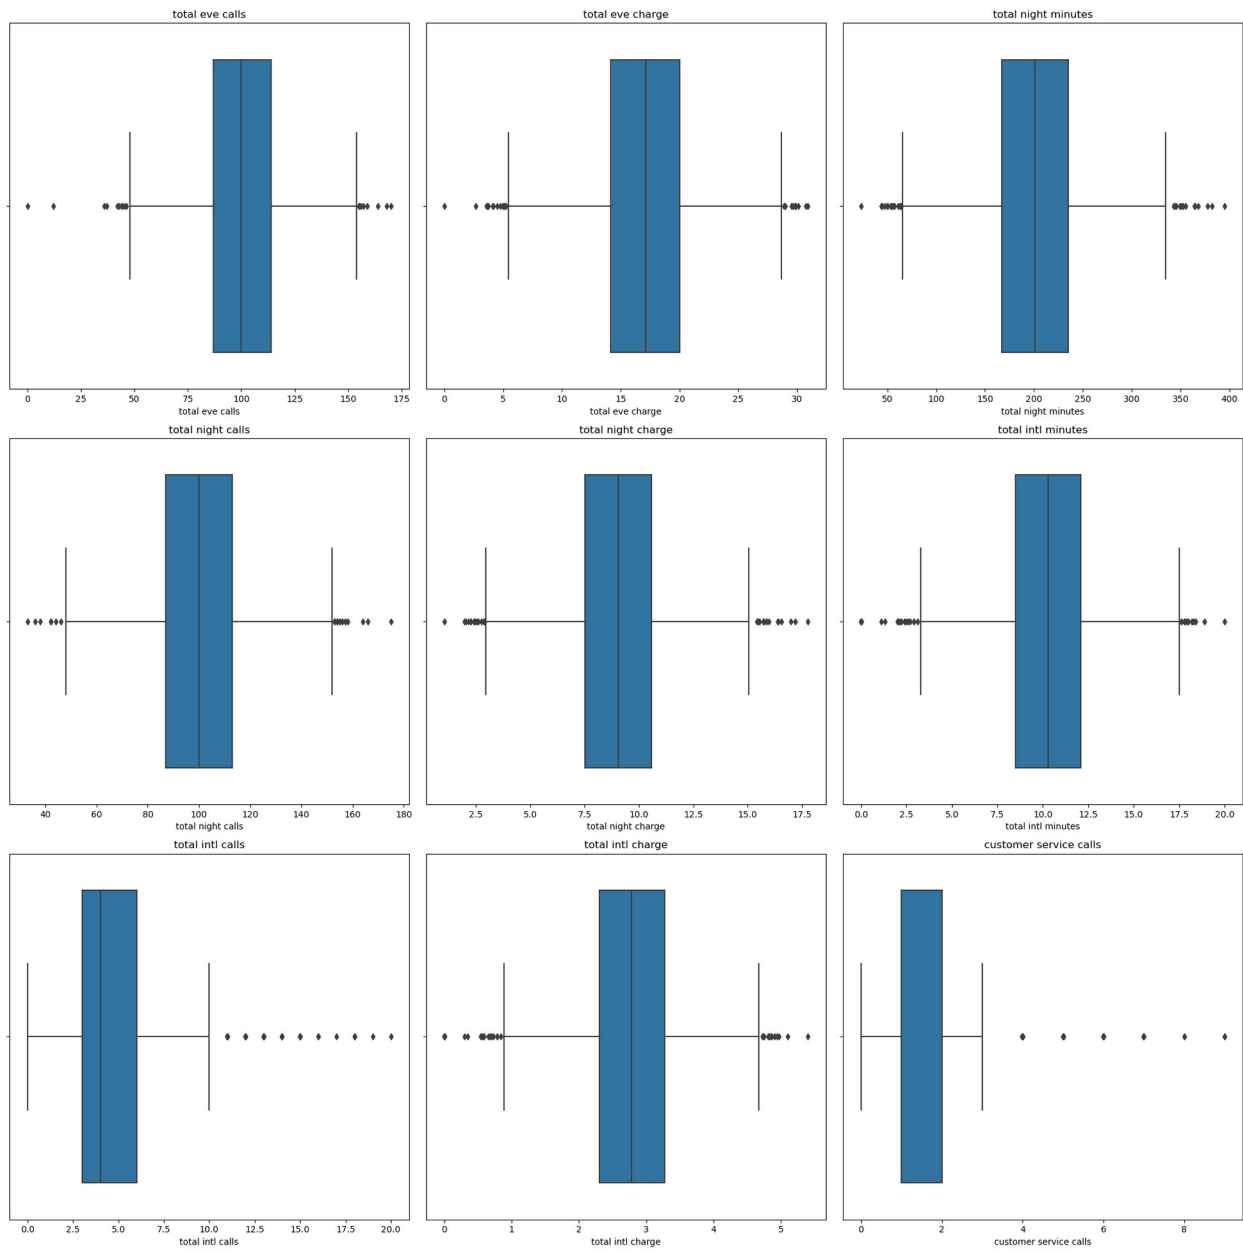
To correct this, various ensemble methods will be applied when modelling.

Visualizing the Outliers

```
cols_boxplot = df.columns[[10, 11, 12, 13, 14, 15, 16, 17, 18,]]
fig, axes = plt.subplots(nrows=3, ncols=3, figsize=(20, 20))
axes = axes.flatten()

for i, column in enumerate(cols_boxplot):
    sns.boxplot(x=df[column], ax=axes[i])
    axes[i].set_title(column)

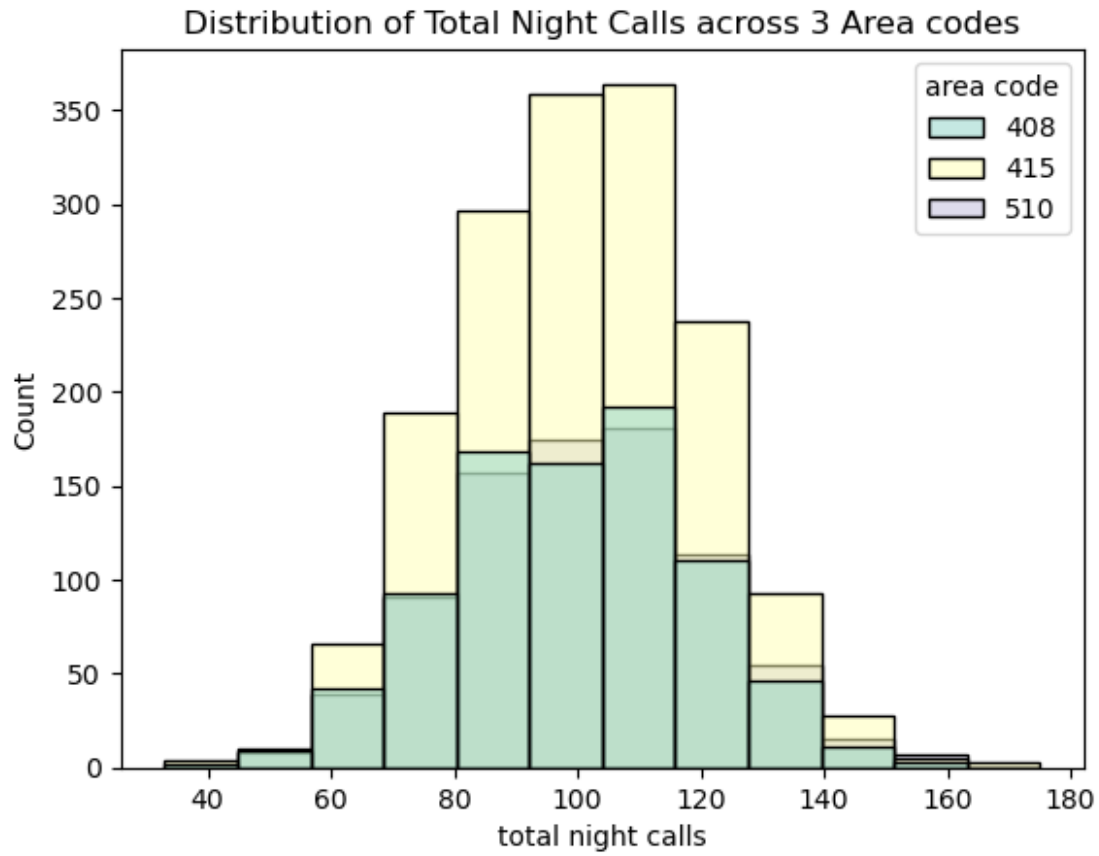
plt.tight_layout()
plt.show()
```



3.2.2 Bivariate Analysis

```
sns.histplot(x="total night calls", kde=False, bins=12, hue= "area
code", data=df, palette="Set3")
plt.title("Distribution of Total Night Calls across 3 Area codes")
plt.show
```


```
<function matplotlib.pyplot.show(close=None, block=None)>
```



Observation:

Area code 415 recorded the highest number of night calls made.

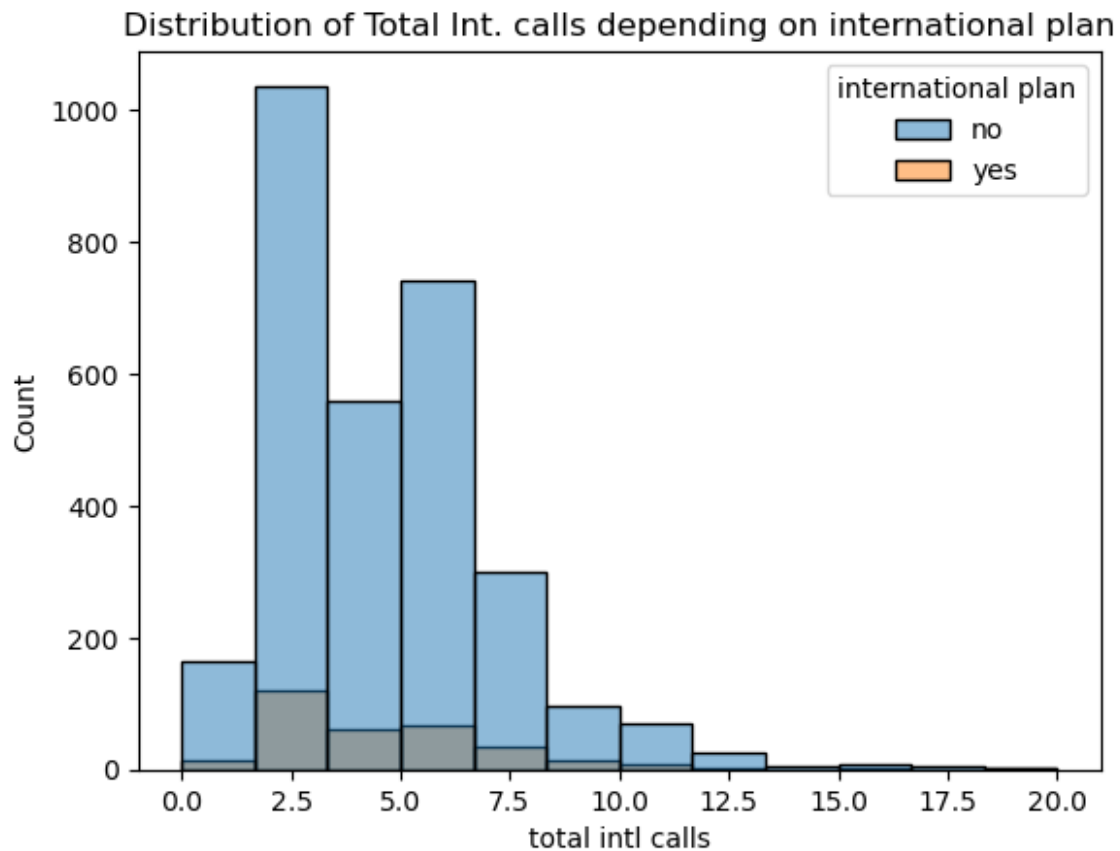
```
sns.histplot(x="total day calls",kde= False, bins=12, hue = "area  
code", data=df, palette= "Set3")  
plt.title("Distribution of Total Day Calls across 3 Area codes")  
plt.show()
```



Observation:

Area code 415 also had the most day calls.

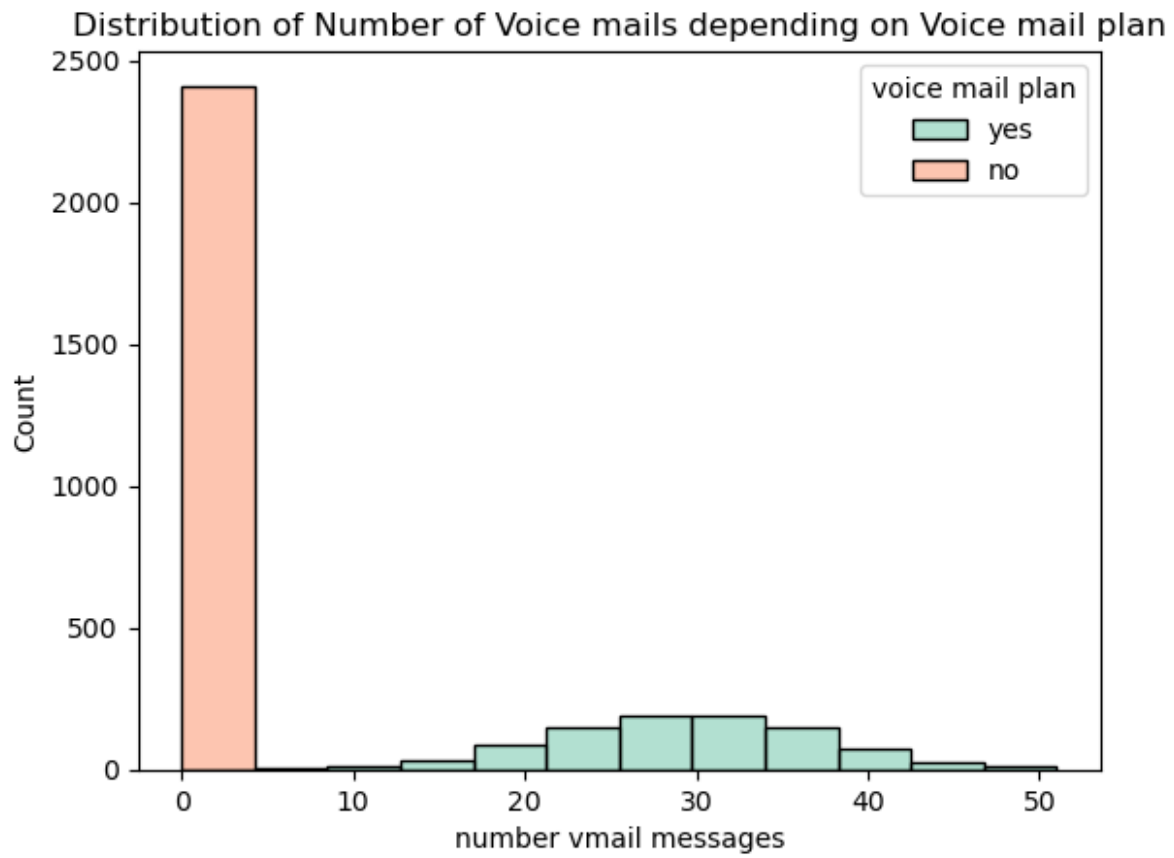
```
sns.histplot(x="total intl calls", data = df, bins=12,  
hue="international plan", kde= False)  
plt.title("Distribution of Total Int. calls depending on international  
plan")  
plt.show()
```



Observation:

Most customers that made international calls did not have an international plan.

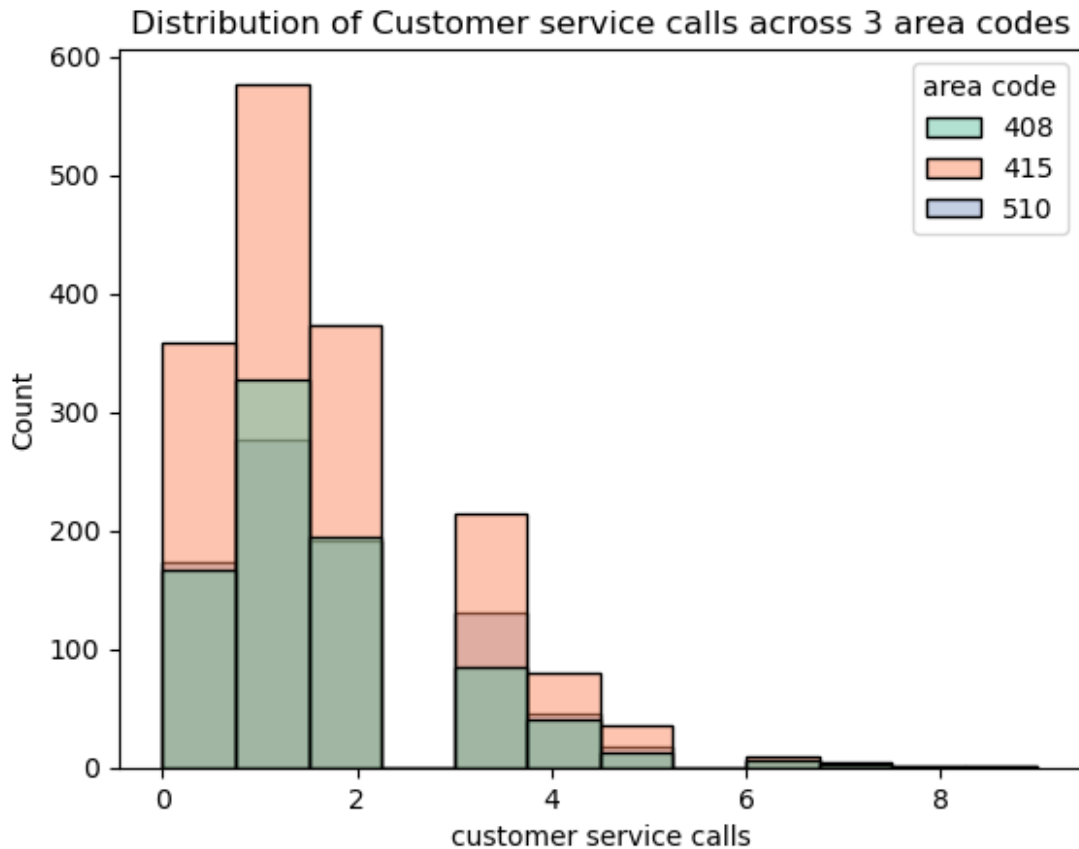
```
sns.histplot(x="number vmail messages", data = df, bins=12, hue="voice  
mail plan", kde= False, palette= "Set2")  
plt.title("Distribution of Number of Voice mails depending on Voice  
mail plan")  
plt.show()
```

Observation:

Many customers who sent voicemail messages had a voice mail plan.

```
sns.histplot(x="customer service calls", bins=12, kde=False,  
hue="area code", data=df, palette="Set2")  
plt.title("Distribution of Customer service calls across 3 area  
codes")  
plt.show()
```



Observation:

Area code 415 once again had the most customer service calls.

It can be therefore confirmed that Area Code 415 was the most active region in using SyriaTel Telecommunications services.

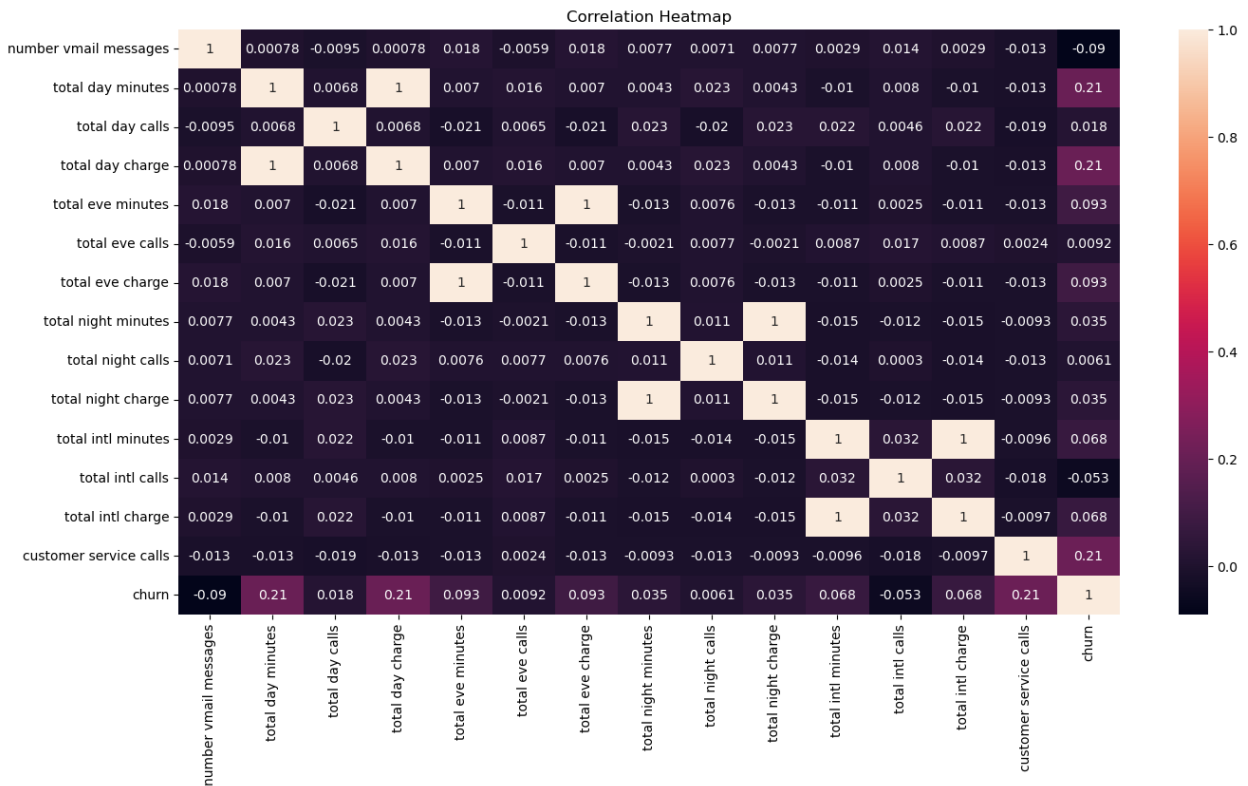
3.2.3 Multivariate Analysis

I will start with a heatmap of all the numerical values to see how they correlate with each other

```
cols_corr=[
    'number vmail messages',
    'total day minutes', 'total day calls', 'total day charge',
    'total eve minutes', 'total eve calls', 'total eve charge',
    'total night minutes', 'total night calls', 'total night
charge',
    'total intl minutes', 'total intl calls', 'total intl charge',
    'customer service calls', 'churn']

plt.figure(figsize=(16,8))
sns.heatmap(data=df[cols_corr].corr(), annot=True)
```

```
plt.title("Correlation Heatmap")
plt.show()
```



4.0 MODELLING

```
# Assigning the target and predictor variables
```

```
X, y = df[df.columns.difference(["churn"]), df["churn"]
```

```
# Mapping categorical columns
```

```
df['churn'] = df['churn'].map({False: 0, True: 1})
```

```
df2 = pd.get_dummies(df, columns=['state', 'area code'])
```

```
df2['international plan'] = df2['international plan'].map({'no': 0, 'yes': 1})
df2['voice mail plan'] = df2['voice mail plan'].map({'no': 0, 'yes': 1})
```

```
# Splitting the dataset
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)
```

```
# Identifying categorical and numerical columns
```

```
categorical_cols = ["state", "international plan", "voice mail plan"]
```



```
# Fitting the pipeline on the training data
pipeline.fit(X_train, y_train)
```

```
# Predicting on the test data
y_pred = pipeline.predict(X_test)
```

```
# Evaluating the model
accuracy = accuracy_score(y_test, y_pred)
print(f'Accuracy: {accuracy:.2f}')
```

Accuracy: 0.86

- The second model improved slightly by 1%
- The third model will entail ensemble methods;

```
# Defining hyperparameter grids for each model
```

```
param_grid_knn = {'n_neighbors': [3, 5, 7, 10]}
```

```
param_grid_svc = {'C': [0.1, 1, 10, 100], 'gamma': [0.01, 0.1, 1,
'scale', 'auto']}
```

```
param_grid_dt = {'max_depth': [None, 5, 10, 20], 'min_samples_split':
[2, 5, 10]}
```

```
param_grid_rf = {'n_estimators': [100, 200, 300], 'max_depth': [None,
5, 10, 20], 'min_samples_split': [2, 5, 10]}
```

```
param_grid_gb = {'n_estimators': [100, 200, 300], 'learning_rate':
[0.01, 0.1, 0.5], 'max_depth': [3, 5, 7]}
```

```
param_grid_xgb = {'n_estimators': [100, 200, 300], 'learning_rate':
[0.01, 0.1, 0.5], 'max_depth': [3, 5, 7]}
```

```
df['churn'] = df['churn'].map({False: 0, True: 1}) #remapping again to
ensure all data is mapped
```

```
# Defining categorical and numerical columns
```

```
categorical_cols = ['state', 'international plan', 'voice mail plan',
'area code']
```

```
numerical_cols = X.select_dtypes(include=['int64',
'float64']).columns.tolist()
```

```
# Creating the column transformer
```

```
preprocessor = ColumnTransformer(
    transformers=[
        ('num', StandardScaler(), numerical_cols),
```

```

        ('cat', OneHotEncoder(drop='first'), categorical_cols)
    ])

# Defining the models and parameter grids
models = {
    'KNN': (KNeighborsClassifier(), {'classifier__n_neighbors': [3, 5, 7, 10]}),
    'SVM': (SVC(), {'classifier__C': [0.1, 1, 10, 100],
                    'classifier__gamma': [0.01, 0.1, 1, 'scale', 'auto']}),
    'Decision Tree': (DecisionTreeClassifier(),
                     {'classifier__max_depth': [None, 5, 10, 20],
                      'classifier__min_samples_split': [2, 5, 10]}),
    'Random Forest': (RandomForestClassifier(),
                     {'classifier__n_estimators': [100, 200, 300], 'classifier__max_depth':
                      [None, 5, 10, 20], 'classifier__min_samples_split': [2, 5, 10]}),
    'Gradient Boosting': (GradientBoostingClassifier(),
                         {'classifier__n_estimators': [100, 200, 300],
                          'classifier__learning_rate': [0.01, 0.1, 0.5],
                          'classifier__max_depth': [3, 5, 7]}),
    'XGBoost': (XGBClassifier(), {'classifier__n_estimators': [100, 200, 300],
                                   'classifier__learning_rate': [0.01, 0.1, 0.5],
                                   'classifier__max_depth': [3, 5, 7]})
}

# Performing RandomizedSearchCV for each model
for name, (model, param_grid) in models.items():
    pipeline = Pipeline(steps=[('preprocessor', preprocessor),
                               ('classifier', model)])
    clf = RandomizedSearchCV(pipeline, param_distributions=param_grid,
                             n_iter=10, cv=5, random_state=42)
    clf.fit(X_train, y_train)
    print(f"Best parameters for {name}: {clf.best_params_}")
    print(f"Training accuracy for {name}: {clf.best_score_}")
    print(f"Test accuracy for {name}: {clf.score(X_test, y_test)}")
    print()

```

```

Best parameters for KNN: {'classifier__n_neighbors': 7}
Training accuracy for KNN: 0.8807246101847361
Test accuracy for KNN: 0.881559220389805

```

```

Best parameters for SVM: {'classifier__gamma': 0.1, 'classifier__C': 10}
Training accuracy for SVM: 0.8998468143713417
Test accuracy for SVM: 0.904047976011994

```

```

Best parameters for Decision Tree: {'classifier__min_samples_split': 5, 'classifier__max_depth': 5}
Training accuracy for Decision Tree: 0.9377342580685962

```

Test accuracy for Decision Tree: 0.9400299850074962

Best parameters for Random Forest: {'classifier__n_estimators': 200, 'classifier__min_samples_split': 5, 'classifier__max_depth': 20}

Training accuracy for Random Forest: 0.9366092571902384

Test accuracy for Random Forest: 0.9400299850074962

Best parameters for Gradient Boosting: {'classifier__n_estimators': 200, 'classifier__max_depth': 5, 'classifier__learning_rate': 0.1}

Training accuracy for Gradient Boosting: 0.9542396582133497

Test accuracy for Gradient Boosting: 0.9505247376311844

Best parameters for XGBoost: {'classifier__n_estimators': 300, 'classifier__max_depth': 7, 'classifier__learning_rate': 0.1}

Training accuracy for XGBoost: 0.9534905945429376

Test accuracy for XGBoost: 0.9535232383808095

5.0 MODEL EVALUATION

- First I trained the first model which was a Logistic Regression as the baseline Model and the second model was the Logistic Regression, but with additional hyperparamaters.
- Based on the Ensemble methods metrics, Gradient Boosting and XGBoost not only performed well on training data but also generalized effectively to unseen test data, suggesting a good balance between underfitting and overfitting.
- Consequently, SyriaTel should prioritize deploying Gradient Boosting and XGBoost models for their churn prediction system.

Next Steps:

Based on the insights derived from the models and feature importance analysis, SyriaTel can implement the following customer retention strategies:

- **Proactive Customer Support:** This could involve regular check-ins or offering dedicated support channels.
- **Personalized Offers and Discounts:** Use predictive models to identify customers at high risk of churn and provide them with personalized offers, discounts, or loyalty programs to enhance their satisfaction and loyalty.
- **Improved Service Plans:** Evaluate and potentially revamp service plans based on the usage patterns identified as critical churn factors. For instance, if customers with low voice mail plan usage are more likely to churn, consider bundling voice mail services with other popular plans.

CONCLUSION

By leveraging the high-performing Gradient Boosting and XGBoost models, SyriaTel can gain valuable insights into customer behavior and implement targeted retention strategies to reduce churn. Continuous monitoring, model updates, and integration with CRM systems will further enhance the effectiveness of these efforts, ultimately contributing to increased customer retention and sustained business growth.